

Computer Systems Security

CSE 466

Fall 2018

Yan Shoshitaishvili

<http://pwn.college>
<http://groups.google.com/group/cse-466>
<https://goo.gl/F2hs9N>

Homework 2 Retrospective

Statistics

88 graded students (+1 auditing)

20 A+ ($\geq 99\%$)

31 A ($\geq 90\%$)

0 B ($\geq 80\%$)

2 C ($\geq 70\%$)


5 D ($\geq 60\%$)

26 E ($< 60\%$)

4 did not attempt

Because of smaller spread (fewer targets leading to many ties), anyone with over 70 flags got an A.

Hw1 Scores and Percents



Top 3

Master hackers:

1. potato - 134 points
2. blub - 124 points
3. smpwnd - 116 points

How did they do it?

Strategies

Lots of evidence of automation (bot accounts, crazy server load, etc).

- Definitely still useful, but less so.

Looks like people have started having pwntools-based "init" scripts.

- `export TERM, SHELL, etc.` before hitting `r.interactive()`

What was the "mastery" hack?

Mastery of the Loading Process

1. File access checks (CAP_DAC_OVERRIDE, then filesystem perms).
2. A new process entry is created.
3. The binary is loaded.
 - a. the rise of Position Independent Executables
4. The libraries are located.
 - a. LD_PRELOAD environment variable, and anything in /etc/ld.so.preload
 - i. (LD_PRELOAD is functionally ignored for setuid binaries)
 - b. DT_RPATH specified in the binary file (can be modified with patchelf)
 - c. LD_LIBRARY_PATH environment variable (can be set in the shell)
 - d. DT_RUNPATH specified in the binary file (can be modified with patchelf)
 - e. system-wide configuration (/etc/ld.so.conf)
 - f. /lib and /usr/lib
5. The libraries are loaded.
 - a. conceptually the same as any other binary (including needing other libraries!)
 - b. relocations updated

Mastery of the Loading Process

1. File access checks (CAP_DAC_OVERRIDE, then filesystem perms).
2. A new process entry is created.
3. The binary is loaded.
 - a. the rise of Position Independent Executables
4. The libraries are located.
 - a. LD_PRELOAD environment variable, and anything in /etc/ld.so.preload
 - i. **(LD_PRELOAD is functionally ignored for setuid binaries)**
 - b. DT_RPATH specified in the binary file (can be modified with patchelf)
 - c. LD_LIBRARY_PATH environment variable (can be set in the shell)
 - d. DT_RUNPATH specified in the binary file (can be modified with patchelf)
 - e. system-wide configuration (/etc/ld.so.conf)
 - f. /lib and /usr/lib
5. The libraries are loaded.
 - a. conceptually the same as any other binary (including needing other libraries!)
 - b. relocations updated

Mastery of the Loading Process

1. File access checks (CAP_DAC_OVERRIDE, then filesystem perms).
2. A new process entry is created.
3. The binary is loaded.
 - a. the rise of Position Independent Executables
4. The libraries are located.
 - a. LD_PRELOAD environment variable, and anything in **/etc/ld.so.preload**
 - i. (LD_PRELOAD is functionally ignored for setuid binaries)
 - b. DT_RPATH specified in the binary file (can be modified with patchelf)
 - c. LD_LIBRARY_PATH environment variable (can be set in the shell)
 - d. DT_RUNPATH specified in the binary file (can be modified with patchelf)
 - e. system-wide configuration (**/etc/ld.so.conf**)
 - f. **/lib** and **/usr/lib**
5. The libraries are loaded.
 - a. conceptually the same as any other binary (including needing other libraries!)
 - b. relocations updated

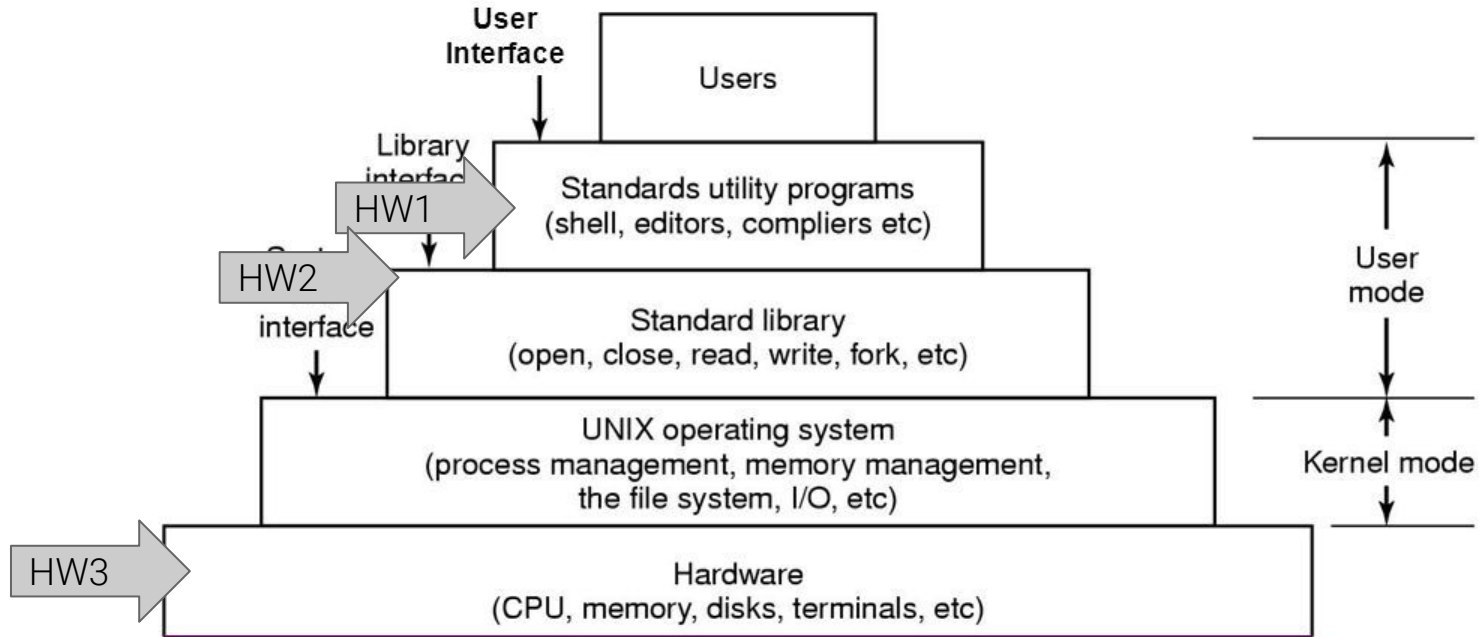
Mastery

Several options to turn *any* arbitrary write into arbitrary execution:

1. Change `/etc/ld.so.conf` to search somewhere where you put a decoy `libc`.
 - a. Need to compile your patched "evil" `libc`.
 - b. Doable, but tedious.
2. Overwrite `/lib/x86_64-linux-gnu/libc.so.6`
 - a. Just as tedious.
3. Change `/etc/ld.so.preload` to preload a simple library that leaks the flag.
 - a. `/etc/ld.so.preload` *DOES* work with SUID binaries.
 - b. Props to: blub, with awesome solves such as **ed**, **tee**, **cp**, and others

CSE 466 Week 3

Binary Brilliance



cat /flag

1. A process is created.
2. Cat is loaded.
 - a. binary
 - b. shared libraries
3. Cat is initialized.
 - a. shared library initializers
 - b. binary initializers
4. **Cat is launched.**
 - a. **__libc_start_main()**
 - b. **main()**
5. **Cat reads its arguments and environment.**
6. **Cat does its thing:**
 - a. **opens /flag**
 - b. **reads the data**
 - c. **writes the data to stdout**
7. Cat terminates.

Software Runs on the CPU

All software eventually runs on your CPU, as low-level machine instructions.









Programs written in compiled languages (C, C++, Rust, etc) run the machine code that they ship as.

Programs written in interpreted languages (Python, JavaScript, etc) are either JIT-compiled immediately before being run, or run by a compiled interpreter (in fancy cases, this interpreter is JIT-compiled itself).






To truly understand how all of this crazy stuff works, we need to understand what happens at the low level!

P.S. Your Computer Organization class should have taught you all of this.


Where do binary files go to?




All our powers combined...



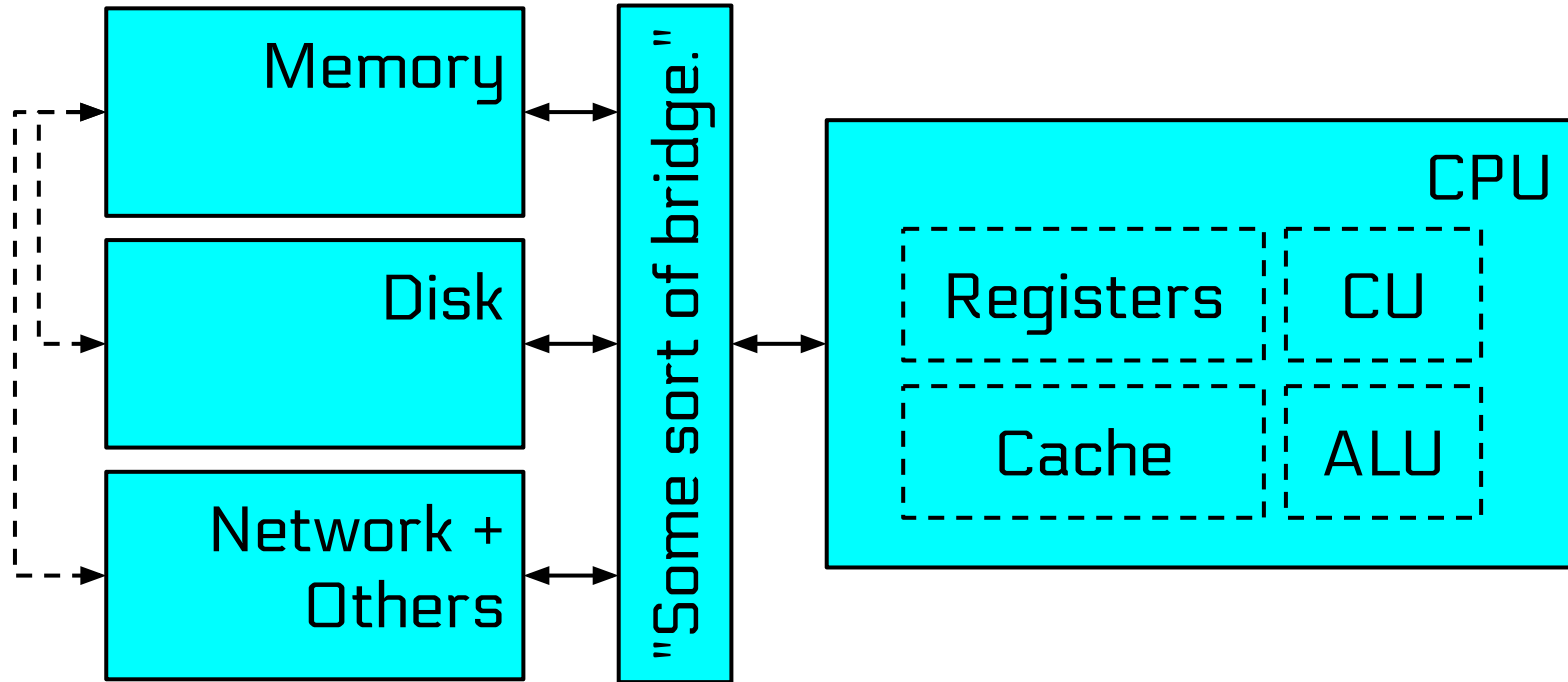
Computer Architecture (at a very high level)



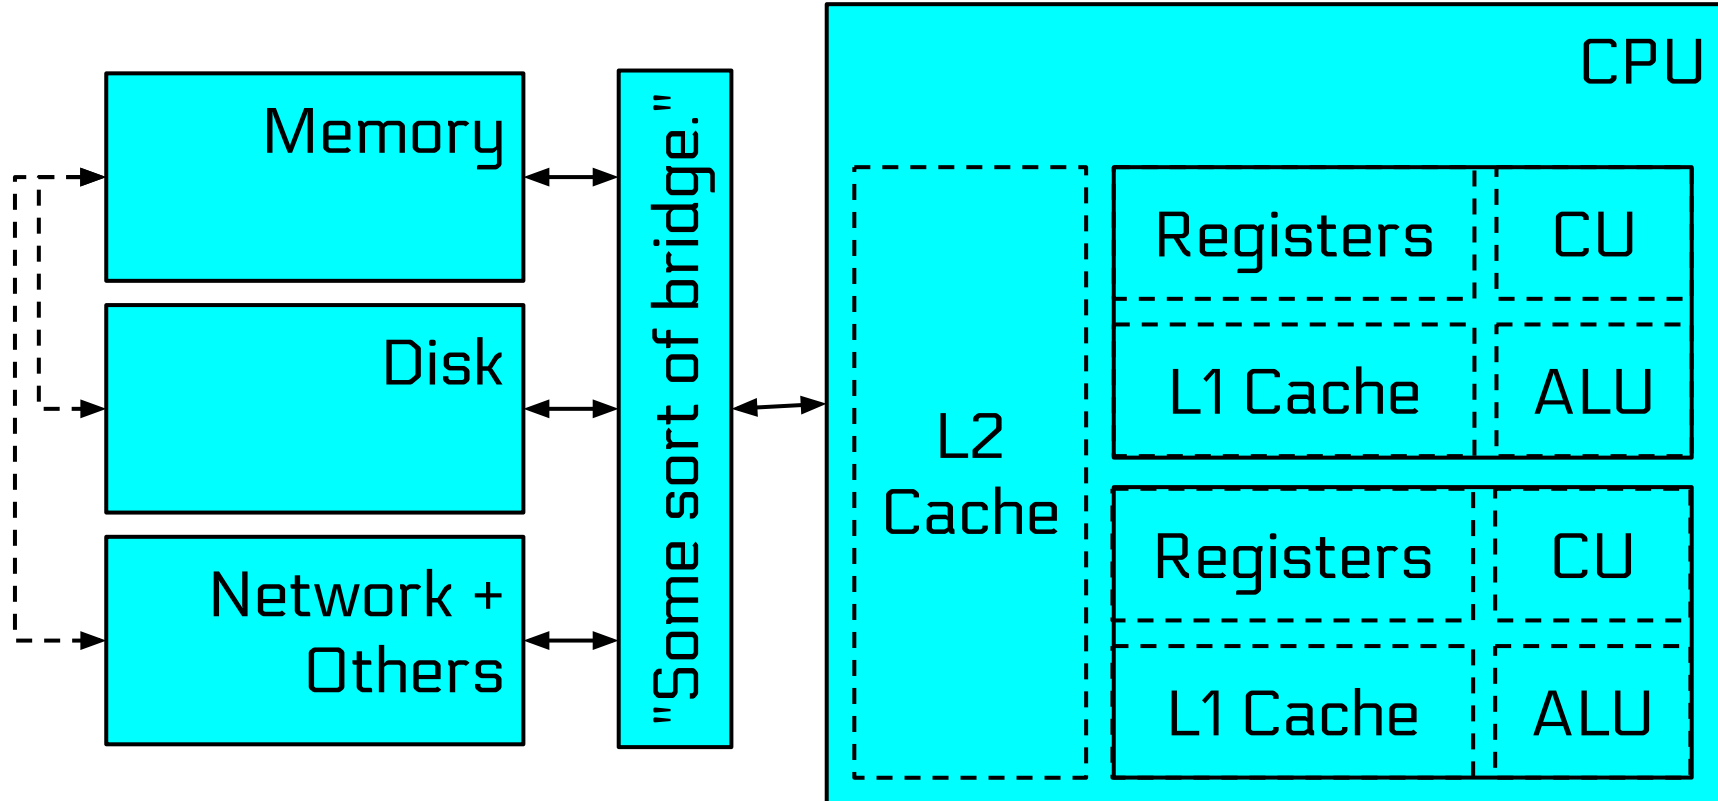
Computer Architecture (drilling down)




Computer Architecture (further down!)



Computer Architecture (as far as we'll go)





John Mauchly (Physicist), John Presper Eckert (Electrical Engineer), John Von Neumann (Mathematician)

John von Neumann, First Draft of a Report on the EDVAC, 1945.

Assembly

The only *true* programming language, as far as a CPU is concerned.

Concepts:

- registers
- memory
- instructions
- stack
- heap

Registers

Registers are very fast, temporary stores for data.

x86: eax, ecx, edx, ebx, **esp, ebp**, esi, edi

amd64: rax, rcx, rdx, rbx, **rsp, rbp**, rsi, rdi, r8, r9, r10, r11, r12, r13, r14, r15

arm: r0, r1, r2, r3, r4, r5, r6, r7, r8, r9, r10, r11, r12, **r13, r14**

The address of the next instruction is in a register:

eip (x86), rip (amd64), r15 (arm)

Various extensions add other registers (x87, MMX, SSE, etc).

Instructions

General form:

OPCODE OPERAND OPERAND, ...

OPCODE - what to do

OPERANDS - what to do it on/with

```
mov rax, rbx
```

```
add rax, 1
```

```
cmp rax, rbx
```

```
jb some_location
```


Instructions (control flow)

Control flow is determined by conditional and unconditional jumps.

Unconditional: call, jmp, ret

Conditional:

je	jump if equal
jne	jump if not equal

jg	jump if greater
jl	jump if less

jle	jump if less than or equal
jge	jump if greater than or equal

ja	jump if above (unsigned)
jb	jump if below (unsigned)

jae	jump if above or equal (unsigned)
jbe	jump if below or equal (unsigned)


js	jump if signed
jns	jump if not signed

jo	jump if overflow
jno	jump if not overflow

jz	jump if zero
jnz	jump if not zero

Instructions (conditionals)

Conditionals key off of the "flags" register: eflags (x86), rflags (amd64), aspr (arm)



Updated by (x86/amd64):

- arithmetic operations
- cmp - subtraction (`cmp rax, rbx`)
- test - and (`test rax, rax`)

je	jump if equal	ZF=1
jne	jump if not equal	ZF=0
jg	jump if greater	ZF=0 and SF=OF
jl	jump if less	SF!=OF
jle	jump if less than or equal	ZF=1 or SF!=OF
jge	jump if greater than or equal	SF=OF
ja	jump if above (unsigned)	CF=0 and ZF=0
jb	jump if below (unsigned)	CF=1
jae	jump if above or equal (unsigned)	CF=0
jbe	jump if below or equal (unsigned)	CF=1 or ZF=1
js	jump if signed	SF=1
jns	jump if not signed	SF=0
jo	jump if overflow	OF=1
jno	jump if not overflow	OF=0
jz	jump if zero	ZF=1
jnz	jump if not zero	ZF=0

Memory (stack)

The stack fulfils four main uses:

1. Track the "callstack" of a program.
 - a. return values are "**pushed**" to the stack during a call and "**popped**" during a ret
2. Contain local variables of functions.
3. Provide scratch space (to alleviate register exhaustion).
4. Pass function arguments (always on x86, only for functions with "many" arguments on other architectures).

Relevant registers (amd64): rsp, rbp

Relevant instructions (amd64): push, pop


Memory (heap)

As we discussed last week, the heap is a libc-managed memory region from which you can allocate (malloc) and deallocate (free) memory.

It doesn't come into play heavily yet, but will feature prominently in future assignments.

Memory (endianess)


Data on most modern systems is stored *backwards*.



Why?

Memory (endianess)

Data on most modern systems is stored *backwards*.



Why?

- Performance (historical)
- Ease of addressing for different sizes.
- (apocryphal) 8086 compatibility

Calling conventions:

Callees and caller functions must agree on argument passing.

Linux x86: push arguments (in reverse order), then call (which pushes return address), return value in eax

Linux amd64: rdi, rsi, rdx, rcx, r8, r9, return value in rax

Linux arm: r0, r1, r2, r3, return value in r0

Registers are *shared* between functions, so calling conventions agree on what registers are protected.

Linux amd64: rbx, rbp, r12, r13, r14, r15 are "callee-saved"

Educational Resources

Rappel (<https://github.com/yrp604/rappel>) lets you explore the effects of instructions.

- easily installable via <https://github.com/zardus/ctf-tools>

pwndevils how2hack:

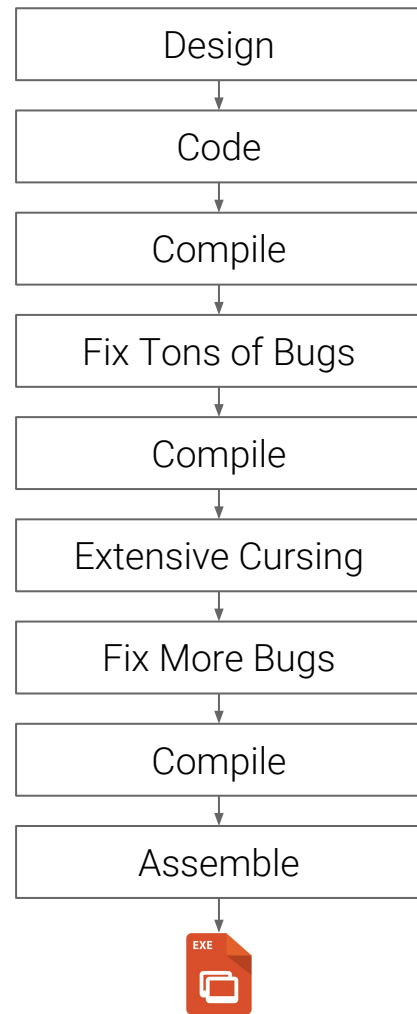
- <http://pwndevils.com/hacking/howtwohack.html>
- web interface to step-by-step understanding of x86
- NOTE: we are using x86_64 today, but many of the concepts are transferable

The Forward Engineering Process

"Forward Engineering" is an overloaded term, but in this context, it is the process of building a program.

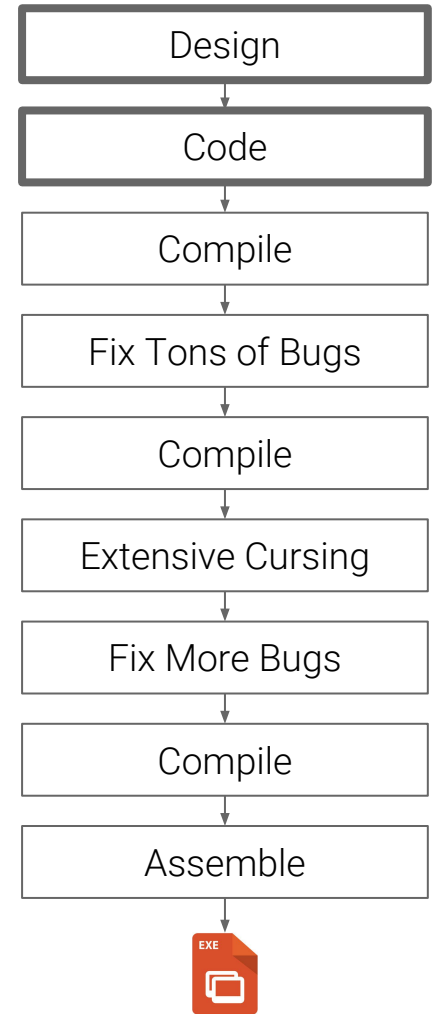
1. Figure out what you want to code.
2. Code it.
3. Compile it (can include JIT).
4. Run it.

At every step, information is lost!



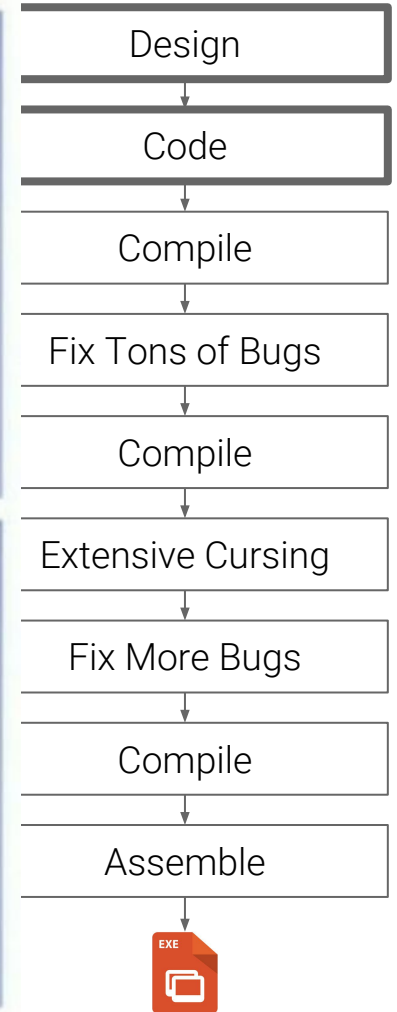
Discussion

What is lost in the transition between Design and Code?



Discussion

What is lost in tl

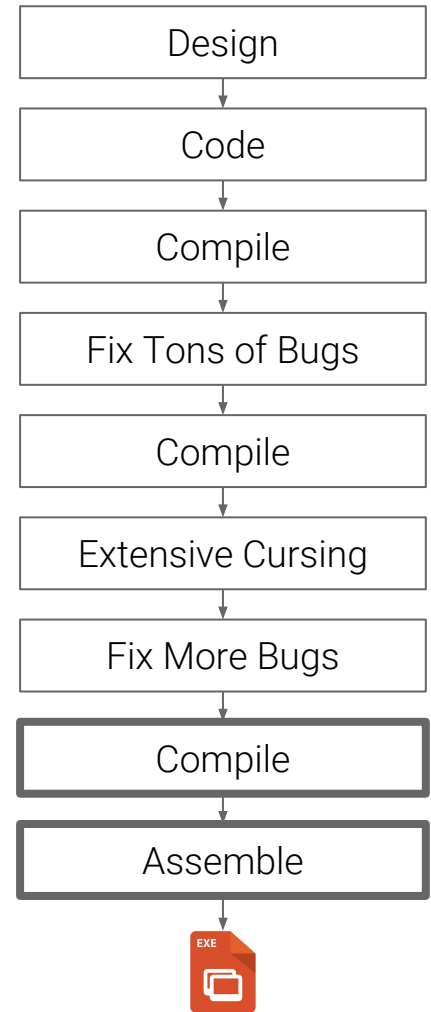


Discussion

What is lost in the compilation process?

- Comments
- Variable names.
- Function names.
- Structure (classes, structs, etc) data.
- Sometimes, entire algorithms (optimization)!

How do we get it back?

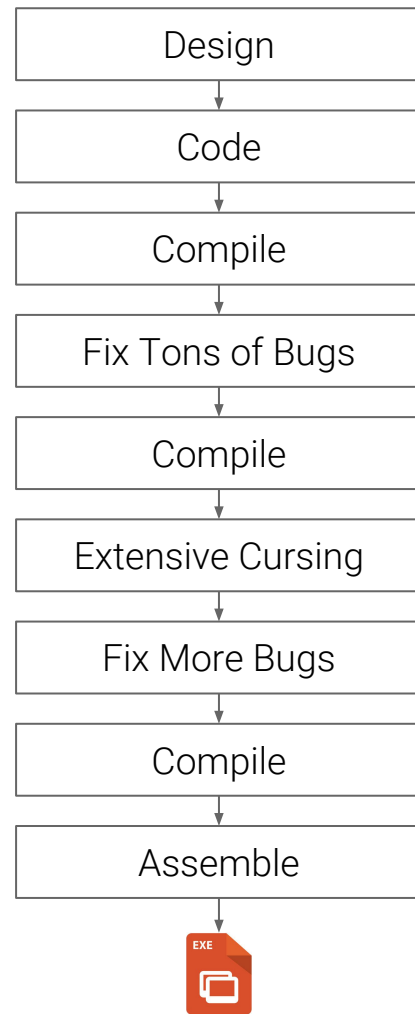


Forward Engineering Tools

Let's look at some tools:

- Your IDE (we will not look at this).
- The GNU C Compiler.
- The GNU Assembler.
 - other assemblers
- The GNU Linker.

Viola! An ELF is born.




The Reverse Engineering Process

Every step in the reverse-engineering process is imperfect and relies on some amount of human help.

The focus of this class: how do we reverse the design from the binary?

Available tools:

- many disassemblers (we'll look at objdump)
- some decompilers (the good ones are \$\$\$\$\$\$)
- debuggers (gdb)
- simple program analysis tools (ltrace, strace)
- heavyweight program analysis tools



The Reverse Engineering Process


Every
and r

The f
the b



Available tools:

- many disassemblers (we'll look at objdump)
- some decompilers (the good ones are \$\$\$)
- debuggers (gdb)
- simple program analysis tools (ltrace, strace)
- heavyweight program analysis tools




Reverse Engineering - Simple Example

Let's look at a simple "crackme" demo.

Crackmes are small programs designed to test a hacker's reverse engineering skillz.

Inspired by real-world license key verification systems.



Crackme: <https://en.wikipedia.org/wiki/Crackme>

Many many crackmes: <https://crackmes.one>

License Checkers

Back in the stone age, before every cave had internet access, software had to be installable without internet access.

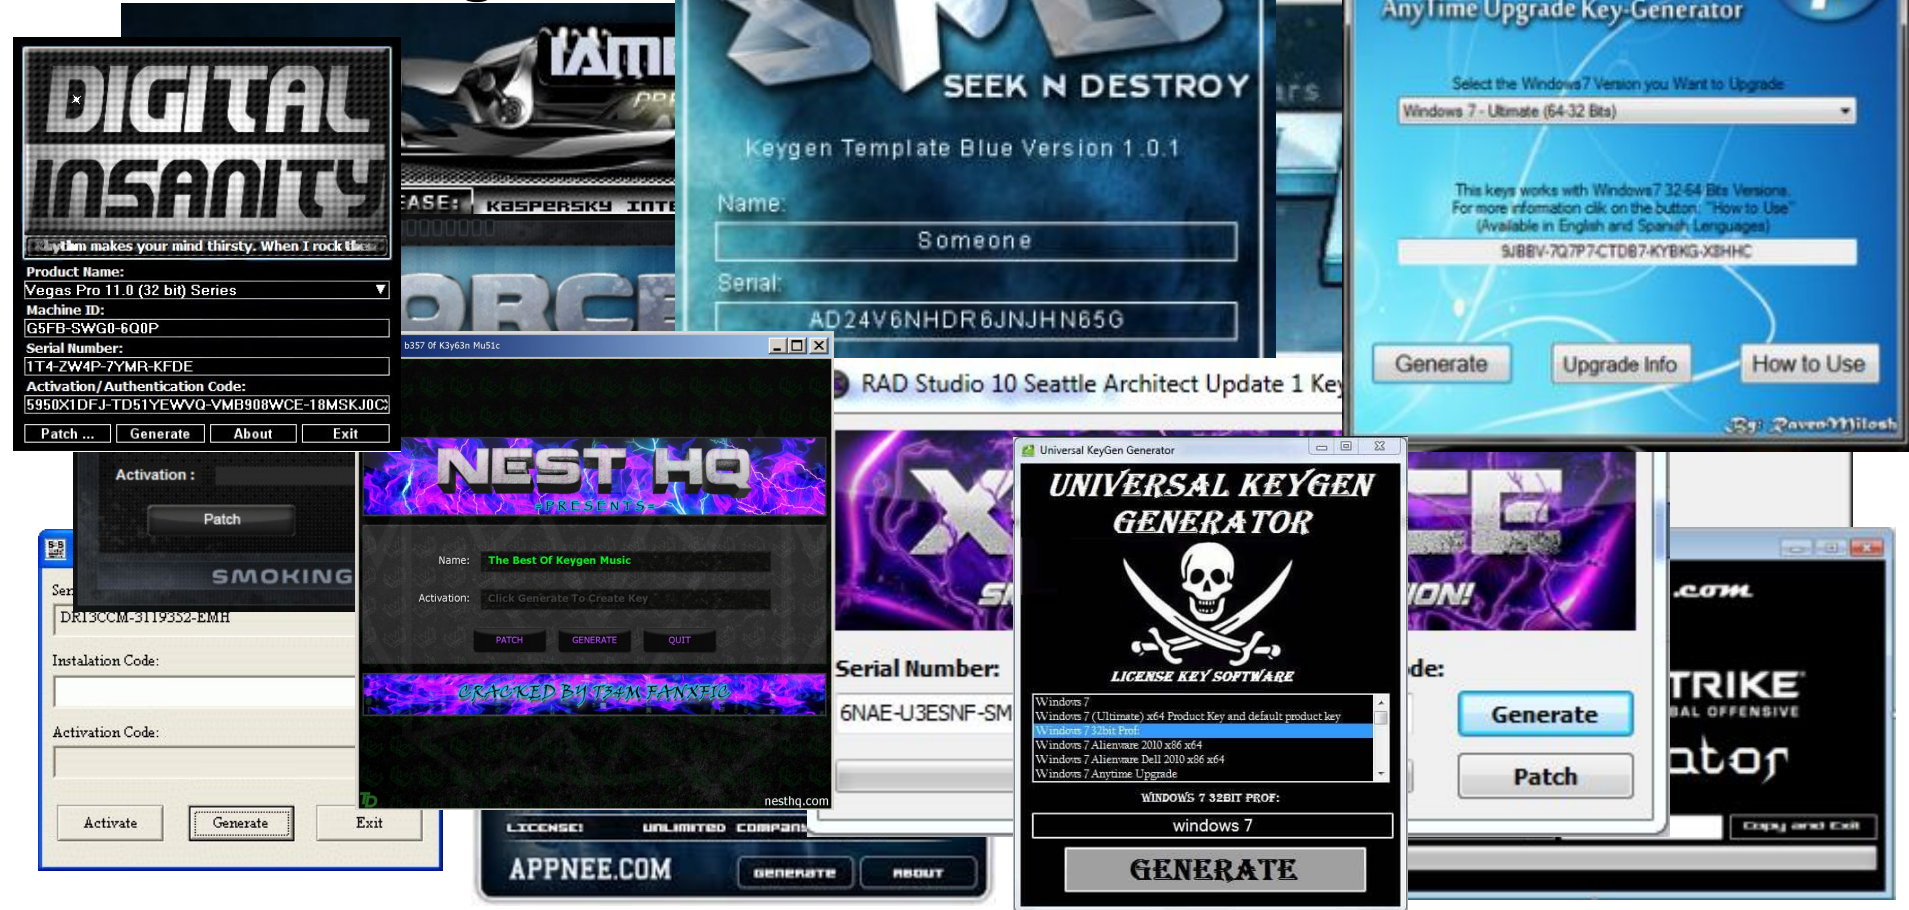
How does the software ensure that it is legitimate?

License key checks!

- Have a secret algorithm that takes in input, performs some calculation, and validates the result.
- Ideally, the company selling the software can generate multiple valid keys.
- Ideally, pirates cannot generate valid keys.

This method implicitly trusts the binary code to keep the secret license key verification algorithm safe!

The Rise of Keygens



The Rise of Keygens



Alternatives to Keygens

1. Cracks that patch the executable to remove the check altogether.
2. Legally purchasing the software.

With ubiquitous internet access (and increased ability to have unbreakable server-side license key checks), keygenning diminished in viability.

But we will explore this lost art in Homework 3!

Homework 3 - Binary Brilliance

Unlike homework 1 and 2, we're diving *in-depth* into the inner workings of binaries, rather than just learning how to use them. It's time to reverse engineer!

The homework:

- netcat to cse466.pwn.college, port 23 (the password will be posted on the class mailing list)
- enter your hacker alias (you can choose a new one) and ASU ID.
- Choose one of your personalized 132 challenges (in /pwn) to attempt. It will be made SUID, and will be able to call `"/get_flag"`.
- If you give that challenge the correct input, it will call `"/get_flag"` and give you the `"/flag"`!
- Log out ("exit" command).
- Provide the flag when asked.
- Scoring is done automatically.
- **2 points** per challenge.

Useful Reversing Tools (and demo)

For simple crackmes, ltrace might be sufficient!

- more info than strace, but still not enough for anything complex

Running the program with multiple different inputs might get you farther.

- ltrace it with input A
- ltrace it again with input B
- see if you can reverse the algorithm from looking at input and output
- still does not scale to complex algorithms

Others:

- gdb (and gdb scripting!)
 - The authority on gdb: <https://sourceware.org/gdb/onlinedocs/gdb/>
- objdump

Useful Reversing Tools (advanced)

This assignment is very doable with ltrace, objdump, and gdb, and I recommend you take that route! That being said, the "magic bullets" for this assignment are advanced program analysis tools.

angr (full disclosure: I am the cofounder of this project)

- advanced program analysis suite that lends itself quite nicely to crackmes
- hard to install (though there is a docker image!), and hard to use
 - you'll probably have to understand how to solve things manually before you can angr them

Dynamic analysis tools:

- <https://github.com/zardus/ctf-tools> has helpful install scripts for tools
- taintgrind will tell you what the program is doing with your input